


Step Forward, Step Backward and Exhaustive Feature Selection | Wrapper Method

- Use combinations of variables to determine predictive power
- Find the best combination of variables
- Computationally expensive than filter method
- Perform better than filter method
- Not recommended on high number of features

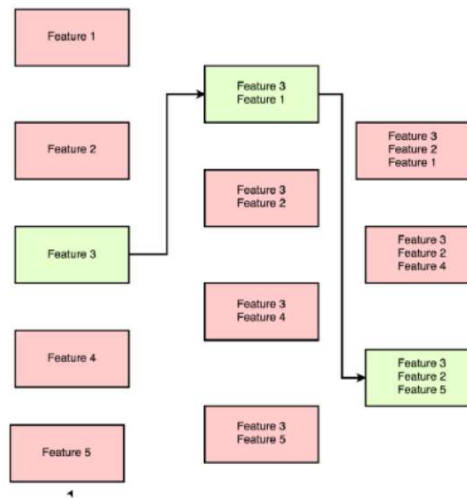
Wrapper Methods

- Type of methods
 - Subset Selection (Exhaustive Feature Selection)
 - Forward Step Selection
 - Backward Step Selection (Recursive Feature Selection)



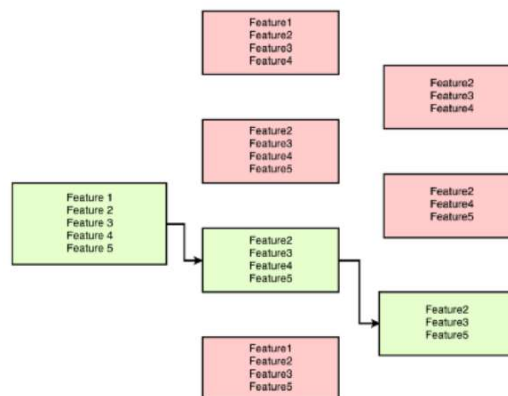
Wrapper Methods

- Forward Step Selection



Wrapper Methods

- Backward Step Selection (Recursive Feature Selection)



Wrapper Methods

- Subset Selection (Exhaustive Feature Selection)
 - fits the model with each possible combinations of N features.
 - requires massive computational power
 - ($Y = B_0$, $Y = B_0 + B_1 \cdot X_1$, $Y = C_0 + C_1 \cdot X_2$, $Y = D_0 + D_1 \cdot X_1 + D_2 \cdot X_2$)
 - Use test error to evaluate model performance

Use of mlxtend in Wrapper Method



In [3]: `!pip install mlxtend`

Requirement already satisfied: mlxtend in c:\users\ibra5\appdata\local\programs\python\python38\lib\site-packages (0.22.0)

Requirement already satisfied: numpy>=1.16.2 in c:\users\ibra5\appdata\roaming\python\python38\site-packages (from mlxtend) (1.23.5)

Requirement already satisfied: pandas>=0.24.2 in c:\users\ibra5\appdata\local\programs\python\python38\lib\site-packages (from mlxtend) (1.5.3)

Requirement already satisfied: matplotlib>=3.0.0 in c:\users\ibra5\appdata\local\programs\python\python38\lib\site-packages (from mlxtend) (3.7.1)

Requirement already satisfied: scipy>=1.2.1 in c:\users\ibra5\appdata\roaming\python\python38\site-packages (from mlxtend) (1.10.1)

Requirement already satisfied: scikit-learn>=1.0.2 in c:\users\ibra5\appdata\roaming\python\python38\site-packages (from mlxtend) (1.2.2)

Requirement already satisfied: joblib>=0.13.2 in c:\users\ibra5\appdata\roaming\python\python38\site-packages (from mlxtend) (1.2.0)

Requirement already satisfied: setuptools in c:\users\ibra5\appdata\local\programs\python\python38\lib\site-packages (from mlxtend) (47.1.0)

Requirement already satisfied: python-dateutil>=2.8.1 in c:\users\ibra5\appdata\local\programs\python\python38\lib\site-packages (from pandas>=0.24.2->mlxtend) (2.8.2)

Requirement already satisfied: pytz>=2020.1 in c:\users\ibra5\appdata\local\programs\python\python38\lib\site-packages (from pandas>=0.24.2->mlxtend) (2022.7.1)

Requirement already satisfied: pyparsing>=2.3.1 in c:\users\ibra5\appdata\local\programs\python\python38\lib\site-packages (from matplotlib>=3.0.0->mlxtend) (3.0.9)

Requirement already satisfied: pillow>=6.2.0 in c:\users\ibra5\appdata\local\programs\python\python38\lib\site-packages (from matplotlib>=3.0.0->mlxtend) (9.4.0)

Requirement already satisfied: cycloper>=0.10 in c:\users\ibra5\appdata\local\programs\python\python38\lib\site-packages (from matplotlib>=3.0.0->mlxtend) (0.11.0)

Requirement already satisfied: importlib-resources>=3.2.0; python_version < "3.10" in c:\users\ibra5\appdata\local\programs\python\python38\lib\site-packages (from matplotlib>=3.0.0->mlxtend) (5.12.0)

Requirement already satisfied: packaging>=20.0 in c:\users\ibra5\appdata\local\programs\python\python38\lib\site-packages (from matplotlib>=3.0.0->mlxtend) (23.0)

Requirement already satisfied: kiwisolver>=1.0.1 in c:\users\ibra5\appdata\local\programs\python\python38\lib\site-packages (from matplotlib>=3.0.0->mlxtend) (1.4.4)

Requirement already satisfied: contourpy>=1.0.1 in c:\users\ibra5\appdata\local\programs\python\python38\lib\site-packages (from matplotlib>=3.0.0->mlxtend) (1.0.7)

Requirement already satisfied: fonttools>=4.22.0 in c:\users\ibra5\appdata\local\programs\python\python38\lib\site-packages (from matplotlib>=3.0.0->mlxtend) (4.39.0)

Requirement already satisfied: threadpoolctl>=2.0.0 in c:\users\ibra5\appdata\roaming\python\python38\site-packages (from scikit-learn>=1.0.2->mlxtend) (3.1.0)

Requirement already satisfied: six>=1.5 in c:\users\ibra5\appdata\local\programs\python\python38\lib\site-packages (from python-dateutil>=2.8.1->pandas>=0.24.2->mlxtend) (1.16.0)

Requirement already satisfied: zipp>=3.1.0; python_version < "3.10" in c:\users\ibra5\appdata\local\programs\python\python38\lib\site-packages (from importlib-resources>=3.2.0; python_version < "3.10"->matplotlib>=3.0.0->mlxtend) (3.15.0)

WARNING: You are using pip version 20.1.1; however, version 23.1.2 is available.

You should consider upgrading via the 'c:\users\ibra5\appdata\local\programs\python\python38\python.exe -m pip install --upgrade pip' command.

How it works

Sequential feature selection algorithms are a family of greedy search algorithms that are used to reduce an initial d-dimensional feature space to a k-dimensional feature subspace where $k < d$.

In a nutshell, SFAs remove or add one feature at the time based on the classifier performance until a feature subset of the desired size k is reached. There are 4 different flavors of SFAs available via the SequentialFeatureSelector:

- Sequential Forward Selection (SFS)
- Sequential Backward Selection (SBS)
- Sequential Forward Floating Selection (SFFS)
- Sequential Backward Floating Selection (SBFS)

Step Forward Selection (SFS)

```
In [4]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [5]: from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor, RandomForestClassifier
from sklearn.metrics import roc_auc_score
from mlxtend.feature_selection import SequentialFeatureSelector as SFS
```

```
In [4]: from sklearn.datasets import load_wine
from sklearn.preprocessing import StandardScaler
```

```
In [19]: data = pd.read_csv('0.9_5subjectslabelled_data.csv', nrows=31437, usecols=lambd
data.head()
```

Out[19]:

	AccX	AccY	AccZ	Gyro_X	Knee Angles	Gait Cycle Phase
0	-0.181472	-0.088708	-0.665352	0.087145	67.223821	5
1	-0.181443	-0.088745	-0.659870	0.103506	67.217858	5
2	-0.183826	-0.089735	-0.654215	0.117635	67.154903	5
3	-0.188545	-0.091706	-0.648504	0.129259	67.011479	5
4	-0.195535	-0.094645	-0.642857	0.138193	66.799616	5

```
In [20]: data.keys()
```

Out[20]: Index(['AccX', 'AccY', 'AccZ', 'Gyro_X', 'Knee Angles', 'Gait Cycle Phase'], dtype='object')

```
In [21]: X = data.drop('Gait Cycle Phase', axis = 1)
y = data['Gait Cycle Phase']

X.shape, y.shape
```

Out[21]: ((31436, 5), (31436,))

```
In [22]: X.isnull().sum()
```

Out[22]:

AccX	0
AccY	0
AccZ	0
Gyro_X	0
Knee Angles	0
dtype: int64	

```
In [23]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, rand
X_train.shape, X_test.shape
```

Out[23]: ((25148, 5), (6288, 5))

Step Forward Feature Selection (SFS)

```
In [24]: sfs = SFS(RandomForestClassifier(n_estimators=100, random_state=0, n_jobs = -1
      k_features = 5,
      forward= True,
      floating = False,
      verbose= 2,
      scoring= 'accuracy',
      cv = 4,
      n_jobs=-1
      ).fit(X_train, y_train)
```

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 2 out of 5 | elapsed: 55.4s remaining: 1.4 min
[Parallel(n_jobs=-1)]: Done 5 out of 5 | elapsed: 57.4s remaining: 0.0s
[Parallel(n_jobs=-1)]: Done 5 out of 5 | elapsed: 57.4s finished

[2023-06-25 02:51:03] Features: 1/5 -- score: 0.37859869572132976[Parallel(n_
jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 4 out of 4 | elapsed: 36.4s remaining: 0.0s
[Parallel(n_jobs=-1)]: Done 4 out of 4 | elapsed: 36.4s finished

[2023-06-25 02:51:39] Features: 2/5 -- score: 0.6335295053284555[Parallel(n_j
obs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 3 out of 3 | elapsed: 21.9s finished

[2023-06-25 02:52:01] Features: 3/5 -- score: 0.8065452521075234[Parallel(n_j
obs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 2 out of 2 | elapsed: 25.9s finished

[2023-06-25 02:52:27] Features: 4/5 -- score: 0.8644425003976459[Parallel(n_j
obs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 1 out of 1 | elapsed: 12.3s finished

[2023-06-25 02:52:39] Features: 5/5 -- score: 0.8954191188166056
```

```
In [26]: sfs.k_feature_names_
```

```
Out[26]: ('AccX', 'AccY', 'AccZ', 'Gyro_X', 'Knee Angles')
```

```
In [27]: sfs.k_feature_idx_
```

```
Out[27]: (0, 1, 2, 3, 4)
```


In [28]: `sfs.k_score_`

Out[28]: 0.8954191188166056

In [29]: `pd.DataFrame.from_dict(sfs.get_metric_dict()).T`

Out[29]:

	feature_idx	cv_scores	avg_score	feature_names	ci_bound	std_dev	std_err
1	(4,)	[0.37967233974868775, 0.3798313981231112, 0.37...	0.378599	(Knee Angles,)	0.002144	0.001337	0.000772
2	(0, 4)	[0.6230316526165103, 0.6433911245427072, 0.637...	0.63353	(AccX, Knee Angles)	0.012398	0.007734	0.004465
3	(0, 2, 4)	[0.8067440750755527, 0.8115158263082551, 0.807...	0.806545	(AccX, AccZ, Knee Angles)	0.006398	0.003991	0.002304
4	(0, 1, 2, 4)	[0.8627326228725942, 0.8719580085891522, 0.860...	0.864443	(AccX, AccY, AccZ, Knee Angles)	0.007046	0.004396	0.002538
5	(0, 1, 2, 3, 4)	[0.8956577063782408, 0.8997932241132496, 0.892...	0.895419	(AccX, AccY, AccZ, Gyro_X, Knee Angles)	0.004436	0.002768	0.001598

```
In [30]: sfs = SFS(RandomForestClassifier(n_estimators=100, random_state=0, n_jobs = -1
      k_features = (1, 5),
      forward= True,
      floating = False,
      verbose= 2,
      scoring= 'accuracy',
      cv = 4,
      n_jobs=-1
      ).fit(X_train, y_train)
```

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 2 out of 5 | elapsed: 48.4s remaining: 1.2 min
[Parallel(n_jobs=-1)]: Done 5 out of 5 | elapsed: 48.9s remaining: 0.0s
[Parallel(n_jobs=-1)]: Done 5 out of 5 | elapsed: 48.9s finished

[2023-06-25 02:54:52] Features: 1/5 -- score: 0.37859869572132976[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 4 out of 4 | elapsed: 32.6s remaining: 0.0s
[Parallel(n_jobs=-1)]: Done 4 out of 4 | elapsed: 32.6s finished

[2023-06-25 02:55:25] Features: 2/5 -- score: 0.6335295053284555[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 3 out of 3 | elapsed: 21.3s finished

[2023-06-25 02:55:46] Features: 3/5 -- score: 0.8065452521075234[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 2 out of 2 | elapsed: 24.3s finished

[2023-06-25 02:56:10] Features: 4/5 -- score: 0.8644425003976459[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 1 out of 1 | elapsed: 11.4s finished

[2023-06-25 02:56:22] Features: 5/5 -- score: 0.8954191188166056
```

```
In [31]: sfs.k_score_
```

```
Out[31]: 0.8954191188166056
```

```
In [32]: sfs.k_feature_names_
```

```
Out[32]: ('AccX', 'AccY', 'AccZ', 'Gyro_X', 'Knee Angles')
```

Step Backward Selection (SBS)

```
In [34]: sfs = SFS(RandomForestClassifier(n_estimators=100, random_state=0, n_jobs = -1
      k_features = (1, 5),
      forward= False,
      floating = False,
      verbose= 2,
      scoring= 'accuracy',
      cv = 4,
      n_jobs=-1
      ).fit(X_train, y_train)
```

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 2 out of 5 | elapsed: 57.2s remaining: 1.4 min
[Parallel(n_jobs=-1)]: Done 5 out of 5 | elapsed: 59.8s remaining: 0.0s
[Parallel(n_jobs=-1)]: Done 5 out of 5 | elapsed: 59.8s finished

[2023-06-25 03:01:28] Features: 4/1 -- score: 0.8644425003976459[Parallel(n_j
obs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 4 out of 4 | elapsed: 30.4s remaining: 0.0s
[Parallel(n_jobs=-1)]: Done 4 out of 4 | elapsed: 30.4s finished

[2023-06-25 03:01:58] Features: 3/1 -- score: 0.8065452521075234[Parallel(n_j
obs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 3 out of 3 | elapsed: 26.2s finished

[2023-06-25 03:02:24] Features: 2/1 -- score: 0.6335295053284555[Parallel(n_j
obs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 2 out of 2 | elapsed: 22.9s finished

[2023-06-25 03:02:47] Features: 1/1 -- score: 0.37859869572132976
```

```
In [37]: sbs = sfs
```

```
In [38]: sbs.k_score_
```

```
Out[38]: 0.8954191188166056
```

```
In [39]: sbs.k_feature_names_
```

```
Out[39]: ('AccX', 'AccY', 'AccZ', 'Gyro_X', 'Knee Angles')
```

Exhaustive Feature Selection (EFS)

```
In [40]: from mlxtend.feature_selection import ExhaustiveFeatureSelector as EFS
```

```
In [42]: efs = EFS(RandomForestClassifier(n_estimators=100, random_state=0, n_jobs=-1),
    min_features= 4,
    max_features= 5,
    scoring='accuracy',
    cv = None,
    n_jobs=-1
    ).fit(X_train, y_train)
```

Features: 6/6

$C(13, 4) + C(13, 5) = 715 + 1287$

```
In [43]: 715 + 1287
```

Out[43]: 2002

```
In [44]: help(efs)
```

Help on ExhaustiveFeatureSelector in module mlxtend.feature_selection.exhaustive_feature_selector object:

```
class ExhaustiveFeatureSelector(sklearn.base.BaseEstimator, sklearn.base.MetaEstimatorMixin)
| ExhaustiveFeatureSelector(estimator, min_features=1, max_features=1, print_progress=True, scoring='accuracy', cv=5, n_jobs=1, pre_dispatch='2*n_jobs', clone_estimator=True, fixed_features=None, feature_groups=None)
```

Exhaustive Feature Selection for Classification and Regression.
(new in v0.4.3)

Parameters

estimator : scikit-learn classifier or regressor

min_features : int (default: 1)
Minimum number of features to select

```
In [45]: efs.best_score_
```

Out[45]: 1.0

In [46]: `efs.best_feature_names_`

Out[46]: ('AccX', 'AccY', 'AccZ', 'Gyro_X')

In [47]: `efs.best_idx_`

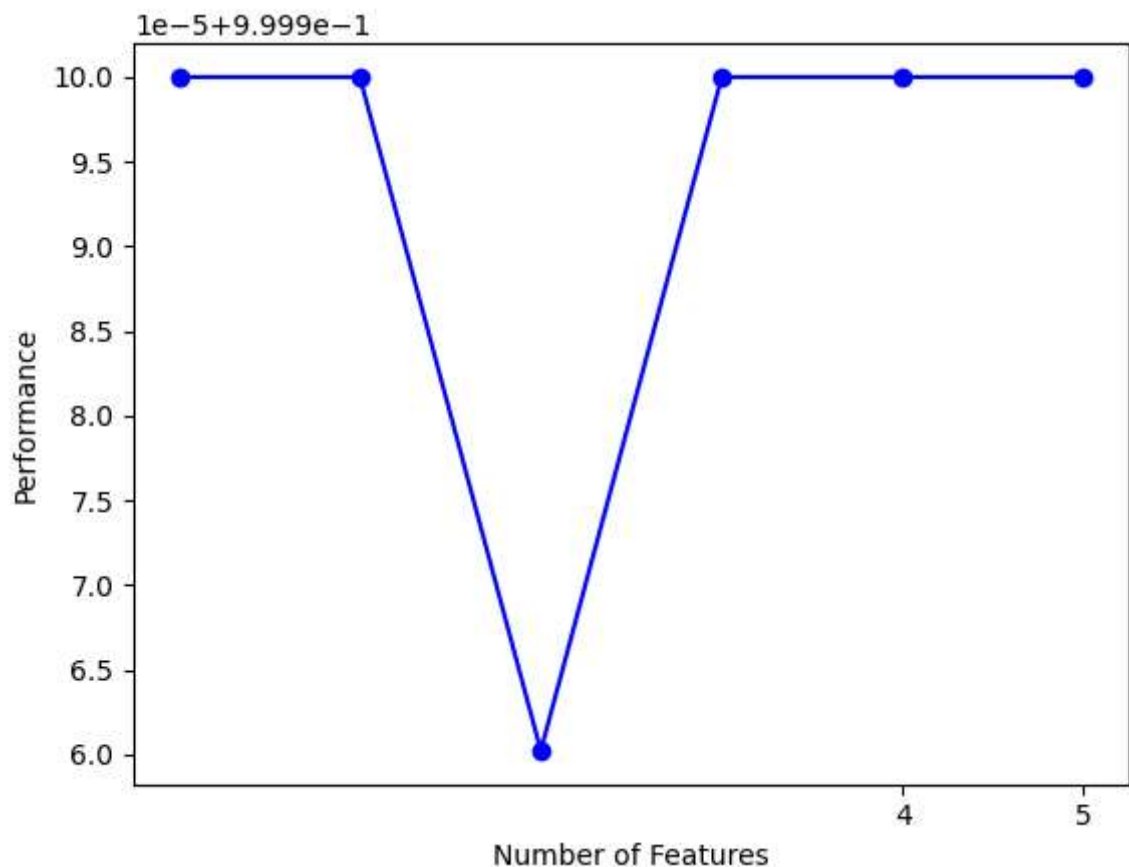
Out[47]: (0, 1, 2, 3)

In [48]: `from mlxtend.plotting import plot_sequential_feature_selection as plot_sfs`

In [49]: `plot_sfs(efs.get_metric_dict(), kind='std_dev')`

C:\Users\ibra5\AppData\Roaming\Python\Python38\site-packages\numpy\core_methods.py:265: RuntimeWarning: Degrees of freedom <= 0 for slice
 ret = _var(a, axis=axis, dtype=dtype, out=out, ddof=ddof,
 C:\Users\ibra5\AppData\Roaming\Python\Python38\site-packages\numpy\core_methods.py:257: RuntimeWarning: invalid value encountered in double_scalars
 ret = ret.dtype.type(ret / rcount)

Out[49]: (<Figure size 640x480 with 1 Axes>,
 <Axes: xlabel='Number of Features', ylabel='Performance'>)



In []:

In []:

In []:	
In []:	
In []:	
In []:	
In []:	
In []:	
In []:	
In []:	
In []:	
In []:	
In []:	
In []:	
In []:	
In []:	
In []:	
In []:	
In []:	
In []:	
In []:	
In []:	
In []:	

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:	
In []:	
In []:	
In []:	
In []:	
In []:	
In []:	
In []:	
In []:	
In []:	
In []:	
In []:	
In []:	
In []:	
In []:	
In []:	
In []:	
In []:	
In []:	